

# Final\_INF 6480

Leslie Garcia

2025-12-13

Load packages

```
library(tidyverse)
```

```
## — Attaching core tidyverse packages ————— tidyverse 2.0.0 —
## ✓ dplyr     1.1.4    ✓ readr     2.1.5
## ✓forcats   1.0.1    ✓ stringr   1.5.2
## ✓ ggplot2   4.0.0    ✓ tibble    3.3.0
## ✓ lubridate 1.9.4    ✓ tidyverse  1.3.1
## ✓ purrr    1.1.0
## — Conflicts ————— tidyverse_conflicts() —
## ✘ dplyr::filter() masks stats::filter()
## ✘ dplyr::lag()   masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
errors
```

```
library(lubridate)
library(broom)
library(reshape2)
```

```
##
## Attaching package: 'reshape2'
##
## The following object is masked from 'package:tidyverse':
##
##     smiths
```

```
library(scales)
```

```
##
## Attaching package: 'scales'
##
## The following object is masked from 'package:purrr':
##
##     discard
##
## The following object is masked from 'package:readr':
##
##     col_factor
```

```
library(dplyr)
```

Read data file

```
border <- read.csv("Border_Crossing_Entry_Data_072025.csv")
```

Filter to California ports data only

```
border_ca <- border %>%
  filter(State == "California")

glimpse(border_ca)
```

```
## Rows: 22,015
## Columns: 10
## $ Port.Name <chr> "Calexico East", "Tecate", "Otay Mesa", "Calexico East", "Ca...
## $ State      <chr> "California", "California", "California", "California", "Cal...
## $ Port.Code   <int> 2507, 2505, 2506, 2507, 2507, 2506, 2505, 2506, 2506, 2582, ...
## $ Border     <chr> "US-Mexico Border", "US-Mexico Border", "US-Mexico Border", ...
## $ Date       <chr> "Jan 2024", "Jan 2024", "Jan 2024", "Jan 2024", "Jan 2024", ...
## $ Measure    <chr> "Personal Vehicle Passengers", "Truck Containers Loaded", "T...
## $ Value      <int> 435768, 5306, 101560, 18199, 247739, 79734, 283, 1815, 2, 20...
## $ Latitude   <dbl> 32.673, 32.576, 32.550, 32.673, 32.673, 32.550, 32.550, 32.5...
## $ Longitude  <dbl> -115.388, -116.626, -116.939, -115.388, -115.388, -116.939, ...
## $ Point      <chr> "POINT (-115.388018 32.673389)", "POINT (-116.626216 32.5764...
```

```
nrow(border_ca) #should be > 0
```

```
## [1] 22015
```

```
table(border_ca[['State']])
```

```
##
## California
##      22015
```

Clean data and create holiday indicators

```
border_ca_clean <- border_ca %>%
mutate (
  Date = my(Date),
  Year = year(Date),
  Month = month(Date),
  Month_Name = month(Date, label = TRUE, abbr = TRUE),

# ORIGINAL holiday definition (July, Nov, Dec)
HolidayMonth = if_else(Month %in% c(7, 11, 12),
                       "Holiday", "Non-Holiday"),
HolidayMonth = factor(HolidayMonth),

# NEW: Spring/Easter months (March & April)
Spring_Easter_Month = Month %in% c(3, 4),

# NEW: Memorial Day month (May)
Memorial_Month = Month == 5,

# OPTIONAL: broader "major travel" months
# (Spring break/Easter, Memorial Day, July 4, Thanksgiving, Christmas/New Year)
Holiday_MajorTravel = if_else(Month %in% c(3, 4, 5, 7, 11, 12),
                               "Holiday", "Non-Holiday"),
Holiday_MajorTravel = factor(Holiday_MajorTravel),

Border = factor(Border),
Measure = factor(Measure),
`Port.Name` = factor(`Port.Name`),
Value_log = log1p(Value),
Lat_scaled = scale(Latitude)[,1],
Lon_scaled = scale(Longitude)[,1]

) %>%
filter(!is.na(Date), !is.na(Value))

summary(border_ca_clean)
```

```

##          Port.Name      State       Port.Code
## Andrade           :3284 Length:22015     Min.   :2502
## Calexico          :3285 Class  :character  1st Qu.:2503
## Calexico East     :4198 Mode   :character  Median  :2505
## Cross Border Xpress: 93                   Mean    :2505
## Otay Mesa          :4042                   3rd Qu.:2506
## San Ysidro         :3492                   Max.    :2582
## Tecate             :3621

##          Border        Date
## US-Mexico Border:22015 Min.   :1996-01-01
##                           1st Qu.:2002-05-01
##                           Median :2008-09-01
##                           Mean   :2009-01-25
##                           3rd Qu.:2015-02-01
##                           Max.   :2025-06-01
##
##          Measure        Value      Latitude
## Pedestrians        : 2210 Min.   :     0 Min.   :32.54
## Personal Vehicle Passengers: 2123 1st Qu.:     0 1st Qu.:32.55
## Personal Vehicles   : 2123 Median  : 1063 Median  :32.58
## Truck Containers Empty   : 1802 Mean    : 140650 Mean    :32.62
## Trucks              : 1802 3rd Qu.: 63606 3rd Qu.:32.67
## Truck Containers Loaded : 1799 Max.    :3755585 Max.    :32.72
## (Other)             :10156

##          Longitude      Point        Year      Month
## Min.   :-117.0 Length:22015     Min.   :1996   Min.   : 1.000
## 1st Qu.:-116.9 Class  :character 1st Qu.:2002   1st Qu.: 3.000
## Median :-116.6 Mode   :character Median :2008   Median : 6.000
## Mean   :-116.0                   Mean   :2009   Mean   : 6.444
## 3rd Qu.:-115.4                   3rd Qu.:2015   3rd Qu.: 9.000
## Max.   :-114.7                   Max.   :2025   Max.   :12.000
##
##          Month_Name    HolidayMonth Spring_Easter_Month Memorial_Month
## Feb     : 1872   Holiday     : 5421 Mode :logical      Mode :logical
## Mar     : 1868   Non-Holiday:16594 FALSE:18281      FALSE:20148
## May     : 1867                   TRUE :3734      TRUE :1867
## Apr     : 1866
## Jun     : 1865
## Jan     : 1862
## (Other):10815

##          Holiday_MajorTravel  Value_log      Lat_scaled      Lon_scaled
## Holiday    :11022   Min.   : 0.000   Min.   :-1.1397   Min.   :-1.1251
## Non-Holiday:10993   1st Qu.: 0.000   1st Qu.:-1.0361   1st Qu.:-1.0222
##                      Median : 6.970   Median :-0.6514   Median :-0.6644
##                      Mean   : 6.463   Mean   : 0.0000   Mean   : 0.0000
##                      3rd Qu.:11.060   3rd Qu.: 0.7838   3rd Qu.: 0.7511
##                      Max.   :15.139   Max.   : 1.4496   Max.   : 1.5057
##

```

Check for months labeled as holiday

```
border_ca_clean %>%
  count(Year, Month, Month_Name, HolidayMonth, Holiday_MajorTravel,
       Spring_Easter_Month, Memorial_Month) %>%
  arrange(Year, Month) %>%
  head(50)
```

##	Year	Month	Month_Name	Holiday	Month_Holiday	Holiday_MajorTravel	Spring_Easter	Month
## 1	1996	1	Jan	Non-Holiday		Non-Holiday		FALSE
## 2	1996	2	Feb	Non-Holiday		Non-Holiday		FALSE
## 3	1996	3	Mar	Non-Holiday		Holiday		TRUE
## 4	1996	4	Apr	Non-Holiday		Holiday		TRUE
## 5	1996	5	May	Non-Holiday		Holiday		FALSE
## 6	1996	6	Jun	Non-Holiday		Non-Holiday		FALSE
## 7	1996	7	Jul	Holiday		Holiday		FALSE
## 8	1996	8	Aug	Non-Holiday		Non-Holiday		FALSE
## 9	1996	9	Sep	Non-Holiday		Non-Holiday		FALSE
## 10	1996	10	Oct	Non-Holiday		Non-Holiday		FALSE
## 11	1996	11	Nov	Holiday		Holiday		FALSE
## 12	1996	12	Dec	Holiday		Holiday		FALSE
## 13	1997	1	Jan	Non-Holiday		Non-Holiday		FALSE
## 14	1997	2	Feb	Non-Holiday		Non-Holiday		FALSE
## 15	1997	3	Mar	Non-Holiday		Holiday		TRUE
## 16	1997	4	Apr	Non-Holiday		Holiday		TRUE
## 17	1997	5	May	Non-Holiday		Holiday		FALSE
## 18	1997	6	Jun	Non-Holiday		Non-Holiday		FALSE
## 19	1997	7	Jul	Holiday		Holiday		FALSE
## 20	1997	8	Aug	Non-Holiday		Non-Holiday		FALSE
## 21	1997	9	Sep	Non-Holiday		Non-Holiday		FALSE
## 22	1997	10	Oct	Non-Holiday		Non-Holiday		FALSE
## 23	1997	11	Nov	Holiday		Holiday		FALSE
## 24	1997	12	Dec	Holiday		Holiday		FALSE
## 25	1998	1	Jan	Non-Holiday		Non-Holiday		FALSE
## 26	1998	2	Feb	Non-Holiday		Non-Holiday		FALSE
## 27	1998	3	Mar	Non-Holiday		Holiday		TRUE
## 28	1998	4	Apr	Non-Holiday		Holiday		TRUE
## 29	1998	5	May	Non-Holiday		Holiday		FALSE
## 30	1998	6	Jun	Non-Holiday		Non-Holiday		FALSE
## 31	1998	7	Jul	Holiday		Holiday		FALSE
## 32	1998	8	Aug	Non-Holiday		Non-Holiday		FALSE
## 33	1998	9	Sep	Non-Holiday		Non-Holiday		FALSE
## 34	1998	10	Oct	Non-Holiday		Non-Holiday		FALSE
## 35	1998	11	Nov	Holiday		Holiday		FALSE
## 36	1998	12	Dec	Holiday		Holiday		FALSE
## 37	1999	1	Jan	Non-Holiday		Non-Holiday		FALSE
## 38	1999	2	Feb	Non-Holiday		Non-Holiday		FALSE
## 39	1999	3	Mar	Non-Holiday		Holiday		TRUE
## 40	1999	4	Apr	Non-Holiday		Holiday		TRUE
## 41	1999	5	May	Non-Holiday		Holiday		FALSE
## 42	1999	6	Jun	Non-Holiday		Non-Holiday		FALSE
## 43	1999	7	Jul	Holiday		Holiday		FALSE
## 44	1999	8	Aug	Non-Holiday		Non-Holiday		FALSE
## 45	1999	9	Sep	Non-Holiday		Non-Holiday		FALSE
## 46	1999	10	Oct	Non-Holiday		Non-Holiday		FALSE
## 47	1999	11	Nov	Holiday		Holiday		FALSE
## 48	1999	12	Dec	Holiday		Holiday		FALSE
## 49	2000	1	Jan	Non-Holiday		Non-Holiday		FALSE
## 50	2000	2	Feb	Non-Holiday		Non-Holiday		FALSE
##			Memorial_Month	n				
## 1				FALSE	72			
## 2				FALSE	72			

```
## 3      FALSE 72
## 4      FALSE 72
## 5      TRUE 72
## 6      FALSE 72
## 7      FALSE 72
## 8      FALSE 72
## 9      FALSE 72
## 10     FALSE 72
## 11     FALSE 72
## 12     FALSE 72
## 13     FALSE 72
## 14     FALSE 72
## 15     FALSE 72
## 16     FALSE 72
## 17      TRUE 72
## 18     FALSE 72
## 19     FALSE 72
## 20     FALSE 72
## 21     FALSE 72
## 22     FALSE 72
## 23     FALSE 72
## 24     FALSE 72
## 25     FALSE 72
## 26     FALSE 72
## 27     FALSE 72
## 28     FALSE 72
## 29      TRUE 72
## 30     FALSE 72
## 31     FALSE 72
## 32     FALSE 72
## 33     FALSE 72
## 34     FALSE 72
## 35     FALSE 72
## 36     FALSE 72
## 37     FALSE 72
## 38     FALSE 72
## 39     FALSE 72
## 40     FALSE 72
## 41      TRUE 72
## 42     FALSE 72
## 43     FALSE 72
## 44     FALSE 72
## 45     FALSE 72
## 46     FALSE 72
## 47     FALSE 72
## 48     FALSE 72
## 49     FALSE 72
## 50     FALSE 72
```

-Check the counts

```
table(border_ca_clean[["HolidayMonth"]])
```

```
##  
##      Holiday Non-Holiday  
##          5421        16594
```

```
table(border_ca_clean[["Holiday_MajorTravel"]])
```

```
##  
##      Holiday Non-Holiday  
##          11022       10993
```

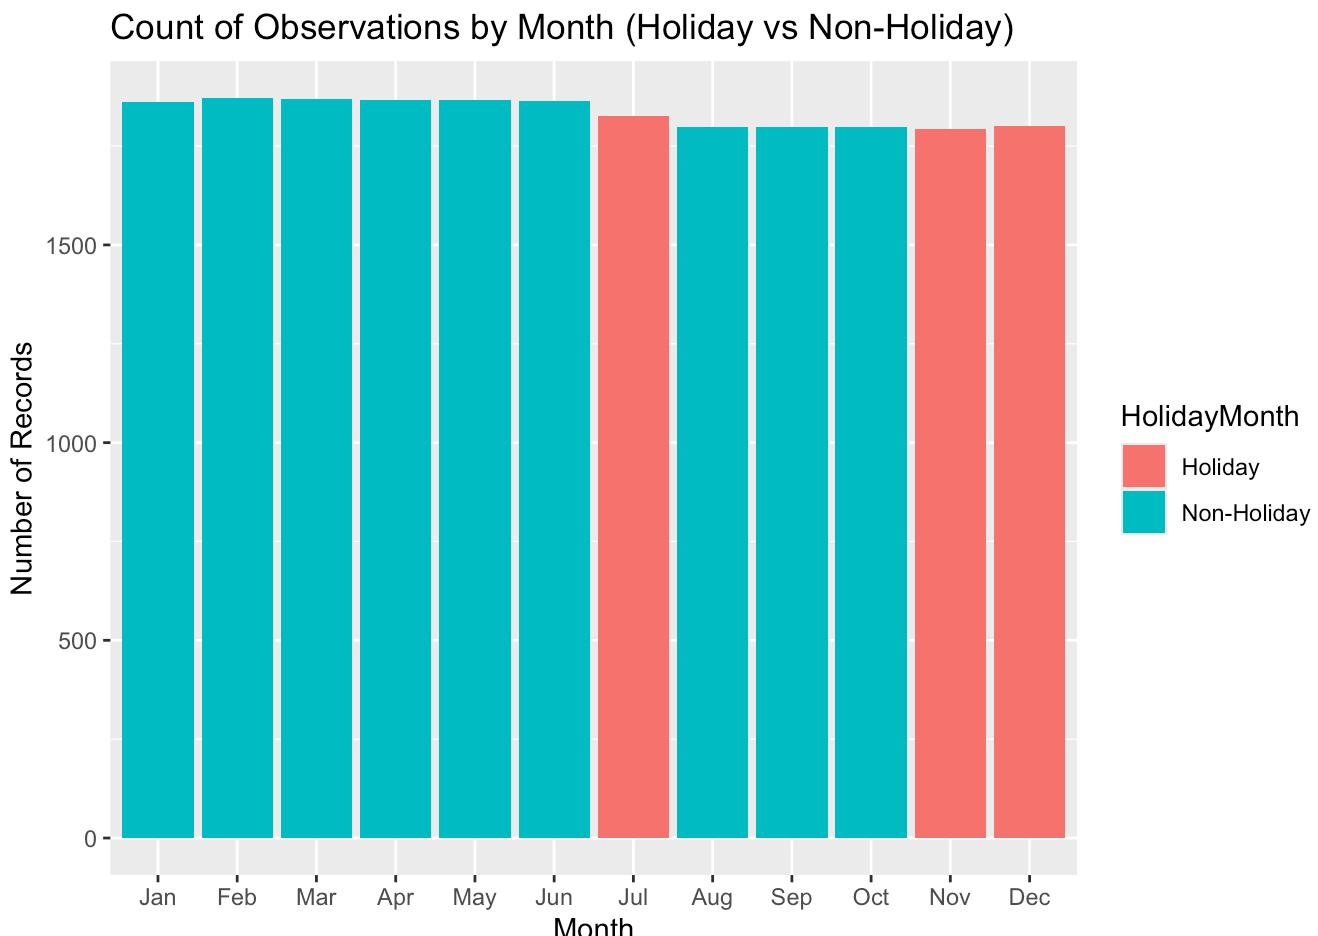
-quick barplot to check labeling

```
table(border_ca_clean[["Holiday_MajorTravel"]])
```

```
##  
##      Holiday Non-Holiday  
##          11022       10993
```

#quick viz to check holiday labeling

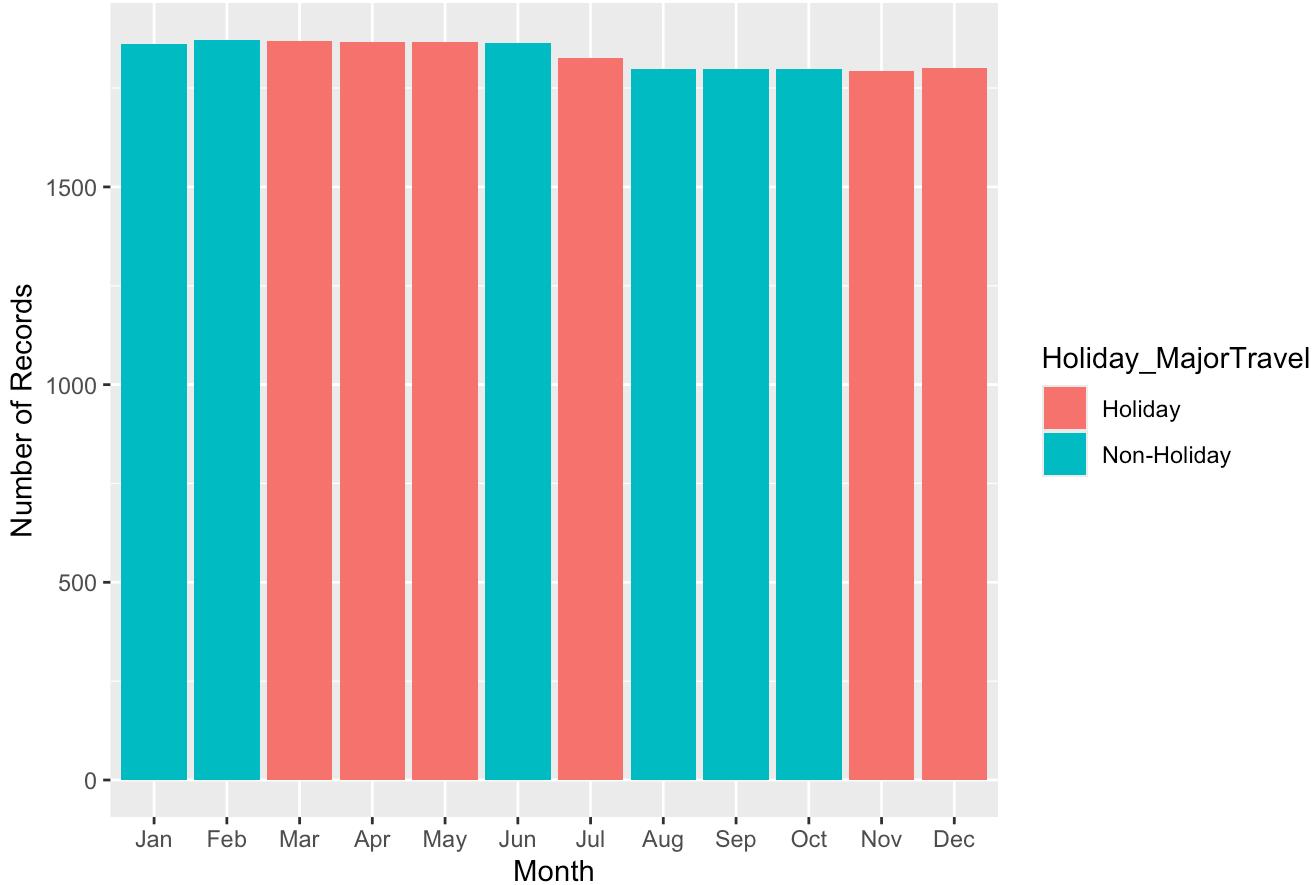
```
ggplot(border_ca_clean, aes(x = Month_Name, fill = HolidayMonth)) +  
  geom_bar() +  
  labs(title = "Count of Observations by Month (Holiday vs Non-Holiday)",  
       x = "Month", y = "Number of Records")
```



Check major travel holiday labels

```
ggplot(border_ca_clean, aes(x = Month_Name, fill = Holiday_MajorTravel)) +  
  geom_bar() +  
  labs(title = "Observations by Month (Major Travel Holiday Definition)",  
       x = "Month", y = "Number of Records")
```

Observations by Month (Major Travel Holiday Definition)



Monthly totals

```
monthly_ca <- border_ca_clean %>%  
  group_by(Date, Month, HolidayMonth) %>%  
  summarise(total_value = sum(Value), .groups = "drop")  
  
max_total <- max(monthly_ca[["total_value"]], na.rm = TRUE)
```

## DESCRIPTIVE STATISTICS

```
#overall summary  
desc_ca <- border_ca_clean %>%  
  summarise(  
    mean_value = mean(Value),  
    median_value = median(Value),  
    sd_value = sd(Value),  
    var_value = var(Value)  
  )  
desc_ca
```

```
##   mean_value median_value sd_value    var_value
## 1 140649.9      1063 369900.6 136826430875
```

```
#Holiday/non-holiday
desc_ca_holiday <- border_ca_clean %>%
  group_by(HolidayMonth) %>%
  summarise(
    mean_value = mean(Value),
    median_value = median(Value),
    sd_value = sd(Value),
    n = n()
  )
desc_ca_holiday
```

```
## # A tibble: 2 × 5
##   HolidayMonth mean_value median_value sd_value     n
##   <fct>          <dbl>        <dbl>     <dbl> <int>
## 1 Holiday        143023.     944    379348.  5421
## 2 Non-Holiday    139875.    1113    366770. 16594
```

```
#Frequency of transportation modes
table(border_ca_clean[["Measure"]])
```

```
##
##           Bus Passengers          Buses
##             1742                  1757
##           Pedestrians Personal Vehicle Passengers
##             2210                  2123
##           Personal Vehicles       Rail Containers Empty
##             2123                  1695
##           Rail Containers Loaded      Train Passengers
##             1594                  1675
##           Trains      Truck Containers Empty
##             1693                  1802
##           Truck Containers Loaded      Trucks
##             1799                  1802
```

```
prop.table(table(border_ca_clean[["Measure"]]))
```

```

##          Bus Passengers           Buses
## 0.07912787 0.07980922
##          Pedestrians Personal Vehicle Passengers
## 0.10038610 0.09643425
##          Personal Vehicles           Rail Containers Empty
## 0.09643425 0.07699296
##          Rail Containers Loaded      Train Passengers
## 0.07240518 0.07608449
##          Trains                   Truck Containers Empty
## 0.07690211 0.08185328
##          Truck Containers Loaded      Trucks
## 0.08171701 0.08185328

```

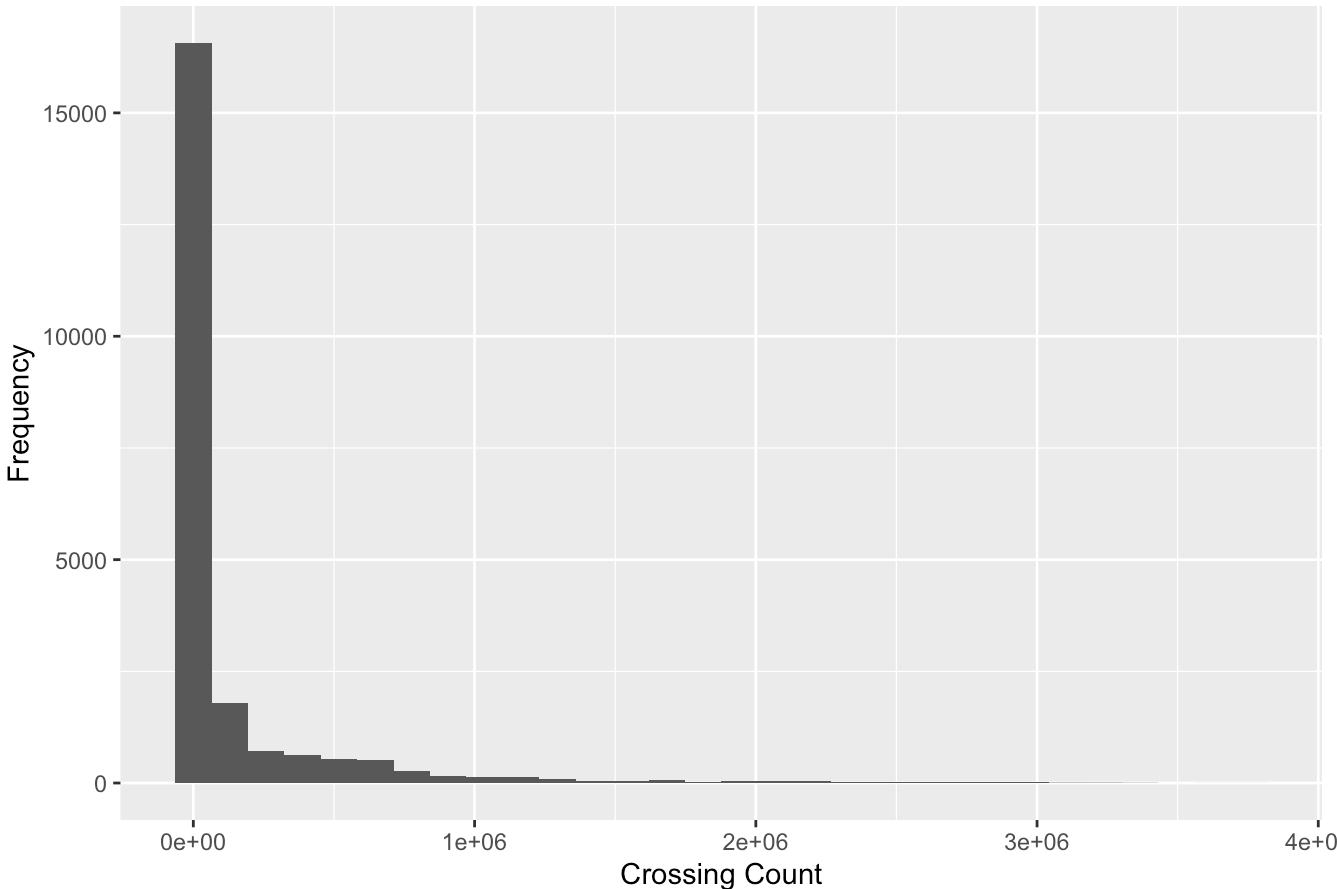
VISUALIZATION - # Histogram of crossing counts

```

ggplot(border_ca_clean, aes(Value)) +
  geom_histogram(bins = 30) +
  labs(title = "Distribution of California Border Crossing Counts",
       x = "Crossing Count", y = "Frequency")

```

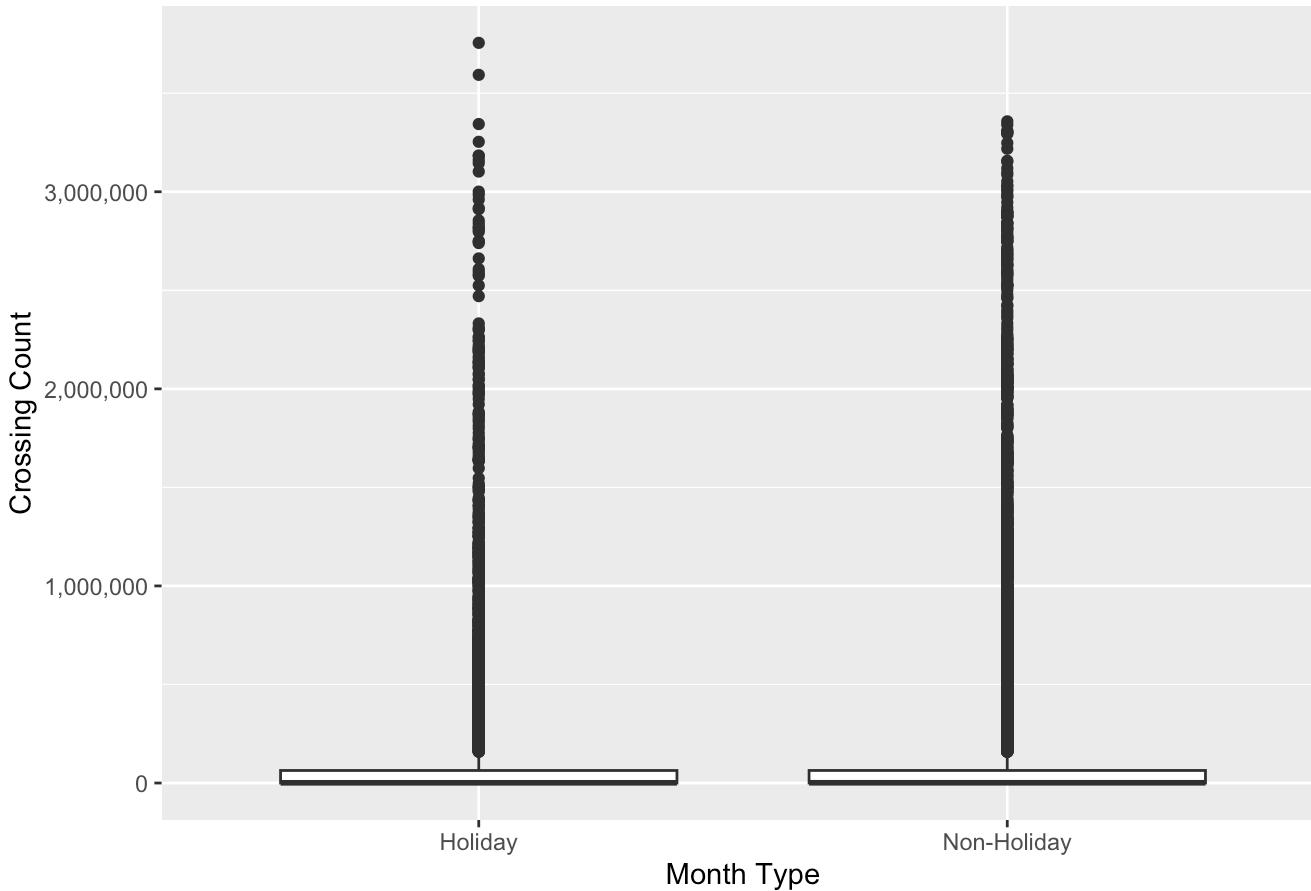
Distribution of California Border Crossing Counts



- Boxplot: Holiday vs Non-Holiday

```
ggplot(border_ca_clean, aes(HolidayMonth, Value)) +
  geom_boxplot() +
  scale_y_continuous(labels = comma) +
  labs(title = "California Border Crossings: Holiday vs Non-Holiday",
       x = "Month Type", y = "Crossing Count")
```

California Border Crossings: Holiday vs Non-Holiday



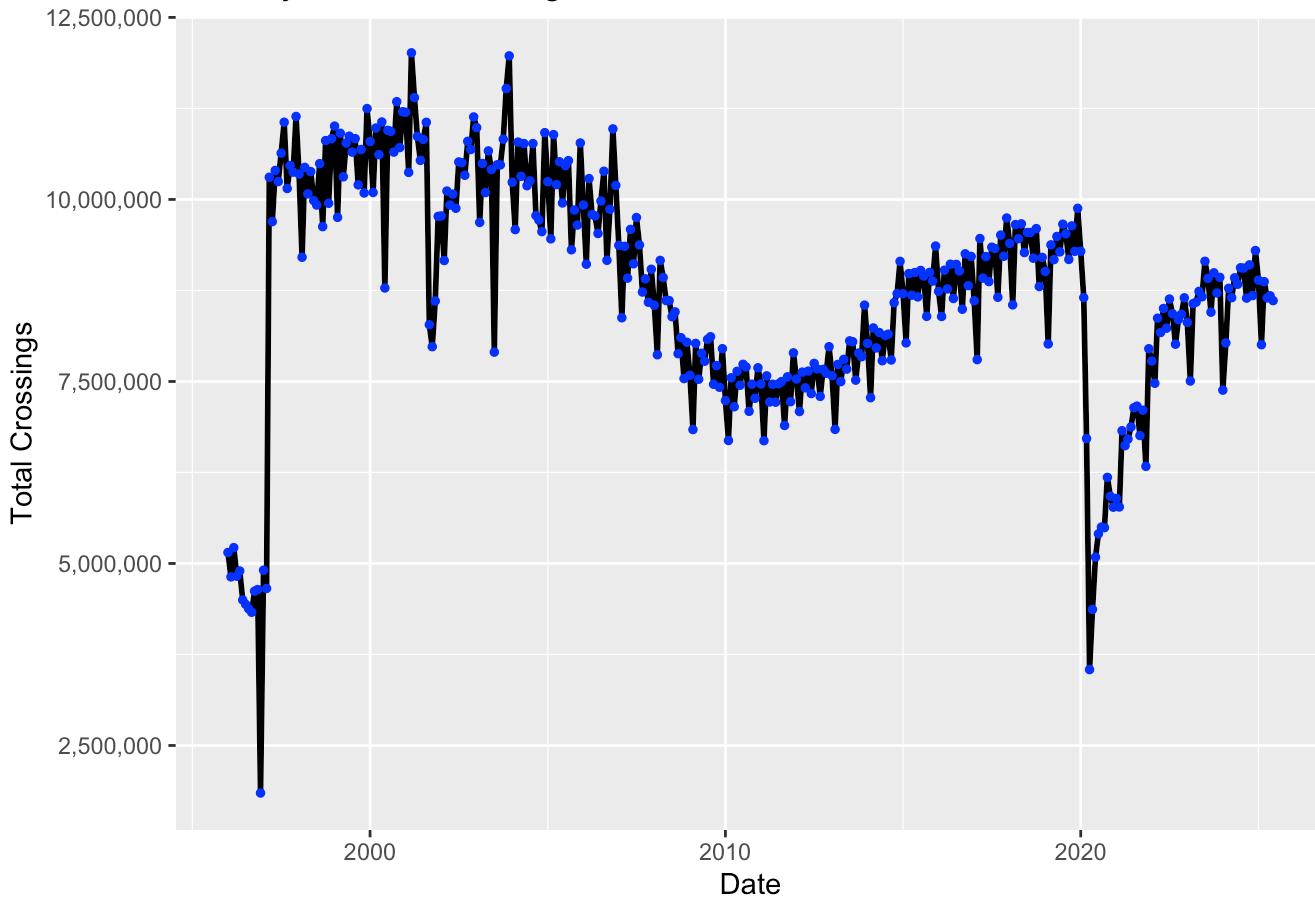
- Time series trend

```
monthly_ca <- border_ca_clean %>%
  group_by(Date) %>%
  summarise(total_value = sum(Value))

ggplot(monthly_ca, aes(Date, total_value)) +
  geom_line(size = 1) +
  geom_point(color = "blue", size = 1) +
  scale_y_continuous(labels = comma) +
  labs(title = "Monthly Border Crossings in California Over Time",
       x = "Date", y = "Total Crossings")
```

```
## Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
## i Please use `linewidth` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```

## Monthly Border Crossings in California Over Time

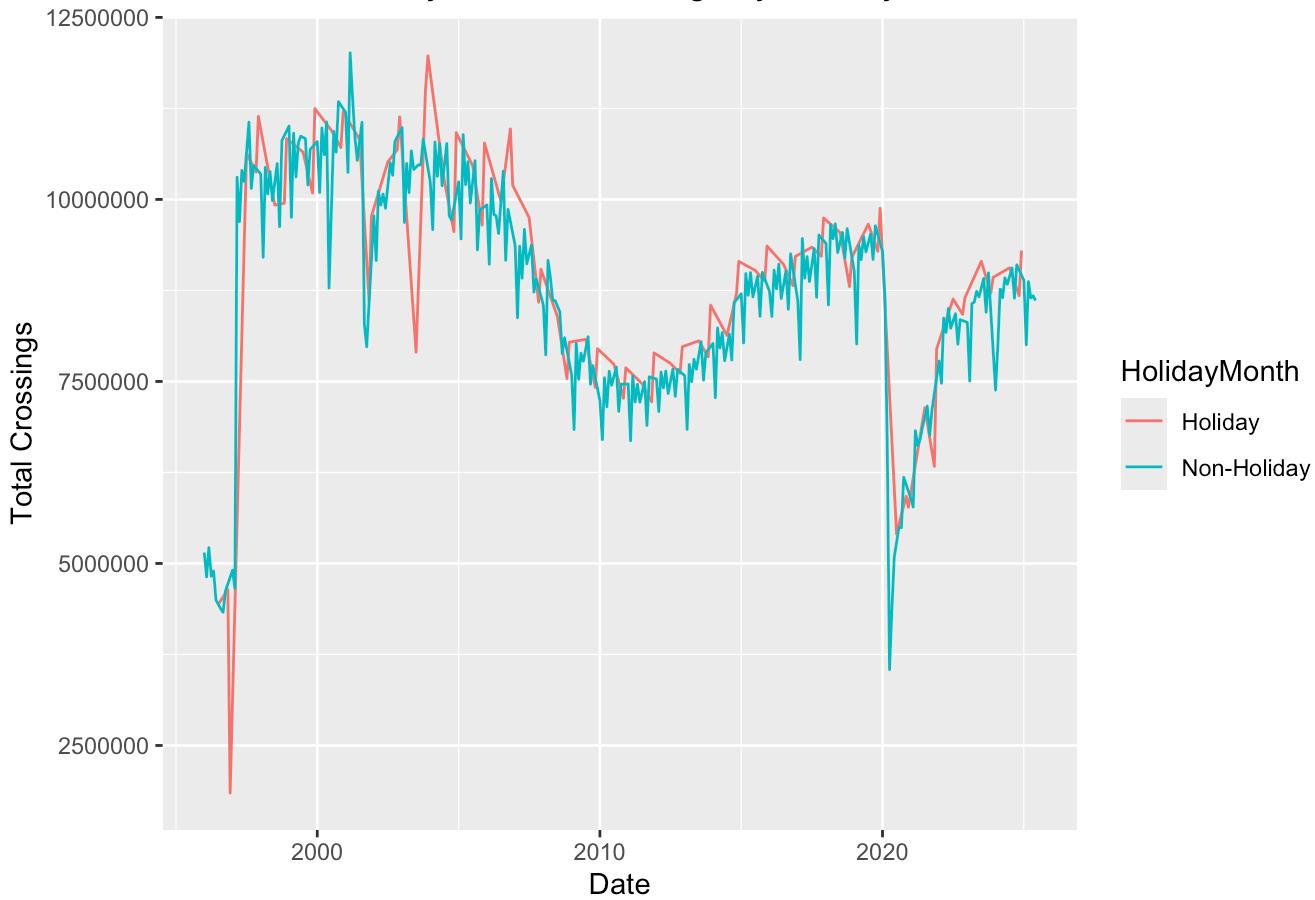


- Trend by holiday

```
monthly_ca_holiday <- border_ca_clean %>%
  group_by(Date, HolidayMonth) %>%
  summarise(total_value = sum(Value), .groups="drop")

ggplot(monthly_ca_holiday,
       aes(Date, total_value, color = HolidayMonth)) +
  geom_line() +
  labs(title = "California Monthly Border Crossings by Holiday Status",
       x = "Date", y = "Total Crossings")
```

## California Monthly Border Crossings by Holiday Status



- Correlation heat map for numeric variables

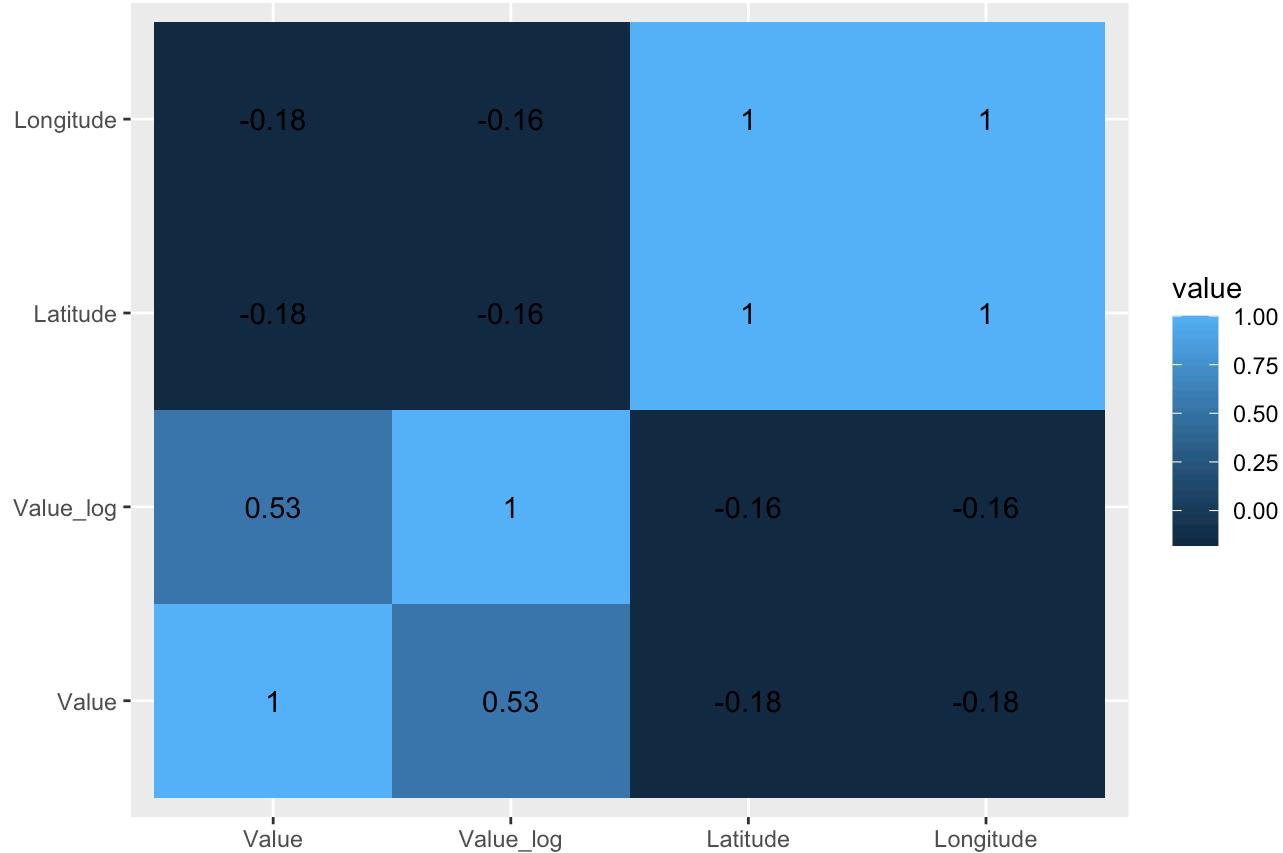
```

num_vars_ca <- border_ca_clean %>%
  select(Value, Value_log, Latitude, Longitude)

cor_mat_ca <- cor(num_vars_ca, use = "complete.obs")
cor_melt_ca <- melt(cor_mat_ca)

ggplot(cor_melt_ca, aes(Var1, Var2, fill = value)) +
  geom_tile() +
  geom_text(aes(label = round(value, 2))) +
  labs(title = "Correlation Heatmap (California Data)",
       x = "", y = "")
  
```

### Correlation Heatmap (California Data)



HYPOTHESIS TESTING - T-test: Are crossings higher during holiday months?

```
ttest_ca <- t.test(Value_log ~ HolidayMonth, data = border_ca_clean)
ttest_ca
```

```
##
## Welch Two Sample t-test
##
## data: Value_log by HolidayMonth
## t = -0.65946, df = 9186.3, p-value = 0.5096
## alternative hypothesis: true difference in means between group Holiday and group Non-Holiday is not equal to 0
## 95 percent confidence interval:
## -0.2080624 0.1033098
## sample estimates:
## mean in group Holiday mean in group Non-Holiday
## 6.423566 6.475942
```

- ANOVA: Differences by transportation mode

```
anova_ca <- aov(Value_log ~ Measure, data = border_ca_clean)
summary(anova_ca)
```

```

##          Df Sum Sq Mean Sq F value Pr(>F)
## Measure      11 363408   33037    3605 <2e-16 ***
## Residuals   22003 201630        9
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

##if significant add
#TukeyHSD(anova_ca)

```

- Chi-square: Is HolidayMonth associated with Measure type?

```

ca_tab <- table(border_ca_clean[["Measure"]], border_ca_clean[["HolidayMonth"]])
chisq.test(ca_tab)

```

```

##
## Pearson's Chi-squared test
##
## data: ca_tab
## X-squared = 0.049906, df = 11, p-value = 1

```

- Correlation: Latitude vs crossing counts

```

cor(border_ca_clean[["Value"]], border_ca_clean[["Latitude"]])

```

```

## [1] -0.1802671

```

## LINEAR REGRESSION MODEL - Modeling dataset

```

model_ca <- border_ca_clean %>%
  select(Value_log, HolidayMonth, Measure,
         Year, Month, Lat_scaled, Lon_scaled) %>%
  na.omit()

nrow(model_ca)

```

```

## [1] 22015

```

- Train-test split

```

set.seed(123)
n <- nrow(model_ca)
train_idx <- sample(seq_len(n), size = 0.7*n)

train_ca <- model_ca[train_idx, ]
test_ca <- model_ca[-train_idx, ]

```

- Fit linear regression model

```

lm_ca <- lm(
  Value_log ~ HolidayMonth + Measure +
  Year + Month + Lat_scaled + Lon_scaled,
  data = train_ca
)

summary(lm_ca)

```

```

##
## Call:
## lm(formula = Value_log ~ HolidayMonth + Measure + Year + Month +
##     Lat_scaled + Lon_scaled, data = train_ca)
##
## Residuals:
##      Min    1Q   Median    3Q   Max 
## -12.8927 -1.5316  0.4565  1.8139  6.1402 
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)                 -99.067100  5.866140 -16.888 < 2e-16 ***
## HolidayMonthNon-Holiday       0.099541  0.065292   1.525   0.127    
## MeasureBuses                  -1.996011  0.114643 -17.411 < 2e-16 ***
## MeasurePedestrians            4.303350  0.108690  39.593 < 2e-16 ***
## MeasurePersonal Vehicle Passengers 5.968575  0.109743  54.387 < 2e-16 ***
## MeasurePersonal Vehicles      5.339536  0.109413  48.802 < 2e-16 *** 
## MeasureRail Containers Empty  -4.157098  0.115982 -35.843 < 2e-16 *** 
## MeasureRail Containers Loaded -5.663742  0.117785 -48.086 < 2e-16 *** 
## MeasureTrain Passengers      -4.900437  0.115503 -42.427 < 2e-16 *** 
## MeasureTrains                 -5.364269  0.114710 -46.764 < 2e-16 *** 
## MeasureTruck Containers Empty -1.181196  0.113169 -10.437 < 2e-16 *** 
## MeasureTruck Containers Loaded -0.788943  0.113701  -6.939 4.11e-12 *** 
## MeasureTrucks                 -0.514230  0.113583  -4.527 6.02e-06 *** 
## Year                          0.052667  0.002920  18.034 < 2e-16 *** 
## Month                         0.007399  0.008192   0.903   0.366    
## Lat_scaled                     14.665729  0.751391  19.518 < 2e-16 *** 
## Lon_scaled                     -15.449658  0.751539 -20.557 < 2e-16 *** 
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.838 on 15393 degrees of freedom
## Multiple R-squared:  0.6865, Adjusted R-squared:  0.6862 
## F-statistic: 2107 on 16 and 15393 DF, p-value: < 2.2e-16

```

- Predictions on test set

```

test_ca[["pred_log"]] <- predict(lm_ca, newdata = test_ca)
head(test_ca)

```

```

##      Value_log HolidayMonth          Measure Year Month Lat_scaled
## 3    11.528415 Non-Holiday     Truck Containers Loaded 2024      1 -1.0360889
## 8     7.504392 Non-Holiday     Truck Containers Empty 2024      1 -0.6514024
## 12   10.821237 Non-Holiday Personal Vehicles 2024      1  1.4495777
## 23   14.516173 Non-Holiday Personal Vehicle Passengers 2024      1 -1.1396584
## 24   8.489205 Non-Holiday           Trucks 2024      1 -0.6514024
## 29  12.969336 Non-Holiday Personal Vehicle Passengers 2024      4  0.7837741
##      Lon_scaled pred_log
## 3    -1.0222204  7.446872
## 8    -0.6643559  7.167444
## 12   1.5056977 10.973995
## 23   -1.1251207 14.275244
## 24   -0.6643559  7.834409
## 29   0.7510953 13.519084

```

- RMSE

```

rmse_ca <- sqrt(mean((test_ca[["pred_log"]]) - test_ca[["Value_log"]])^2))
rmse_ca

```

```

## [1] 2.867307

```

- Test R<sup>2</sup>

```

sst <- sum((test_ca[["Value_log"]]) - mean(test_ca[["Value_log"]]))^2)
sse <- sum((test_ca[["Value_log"]]) - test_ca[["pred_log"]])^2)
r2_test_ca <- 1 - sse/sst
r2_test_ca

```

```

## [1] 0.6799512

```

- Predicted vs Actual Plot

```

ggplot(test_ca, aes(Value_log, pred_log)) +
  geom_point(alpha = 0.4) +
  geom_abline(slope=1, intercept=0, linetype="dashed") +
  labs(title = "Predicted vs Actual (California Border Crossings)",
       x = "Actual log(1 + Value)", y = "Predicted log(1 + Value)")

```

## Predicted vs Actual (California Border Crossings)

